

Psychoinformatics & Neuroinformatics



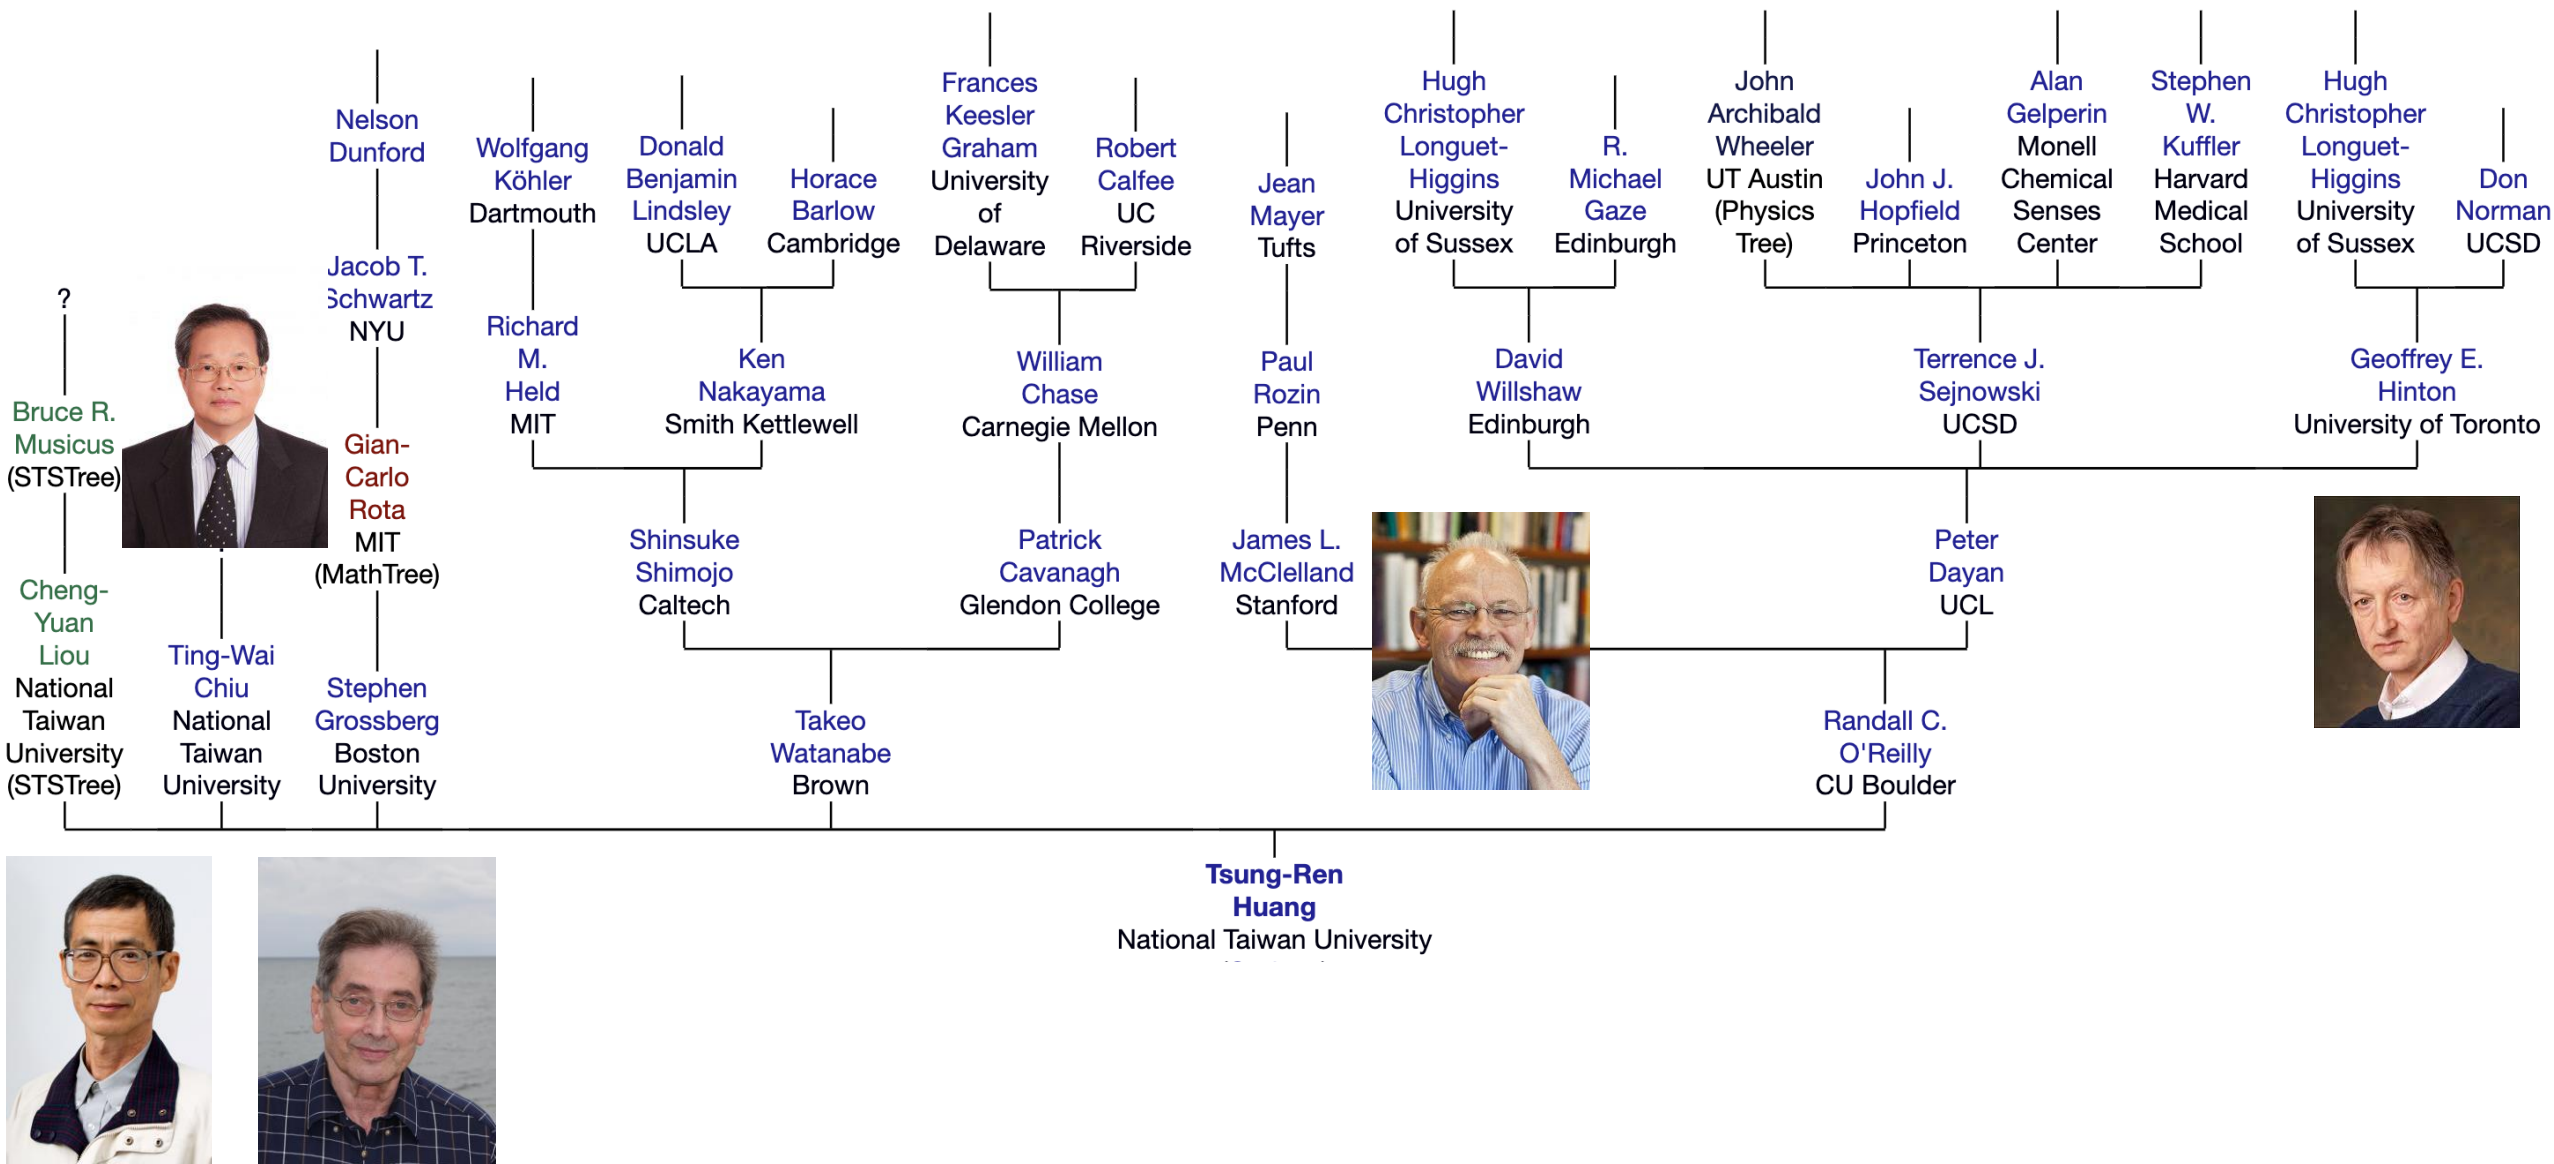
Week 11

Machine Learning (3/3)



by Tsung-Ren (Tren) Huang 黃從仁


Tren's Academic Tree



Artificial Intelligence

Machine Learning

(Deep)
Neural Net

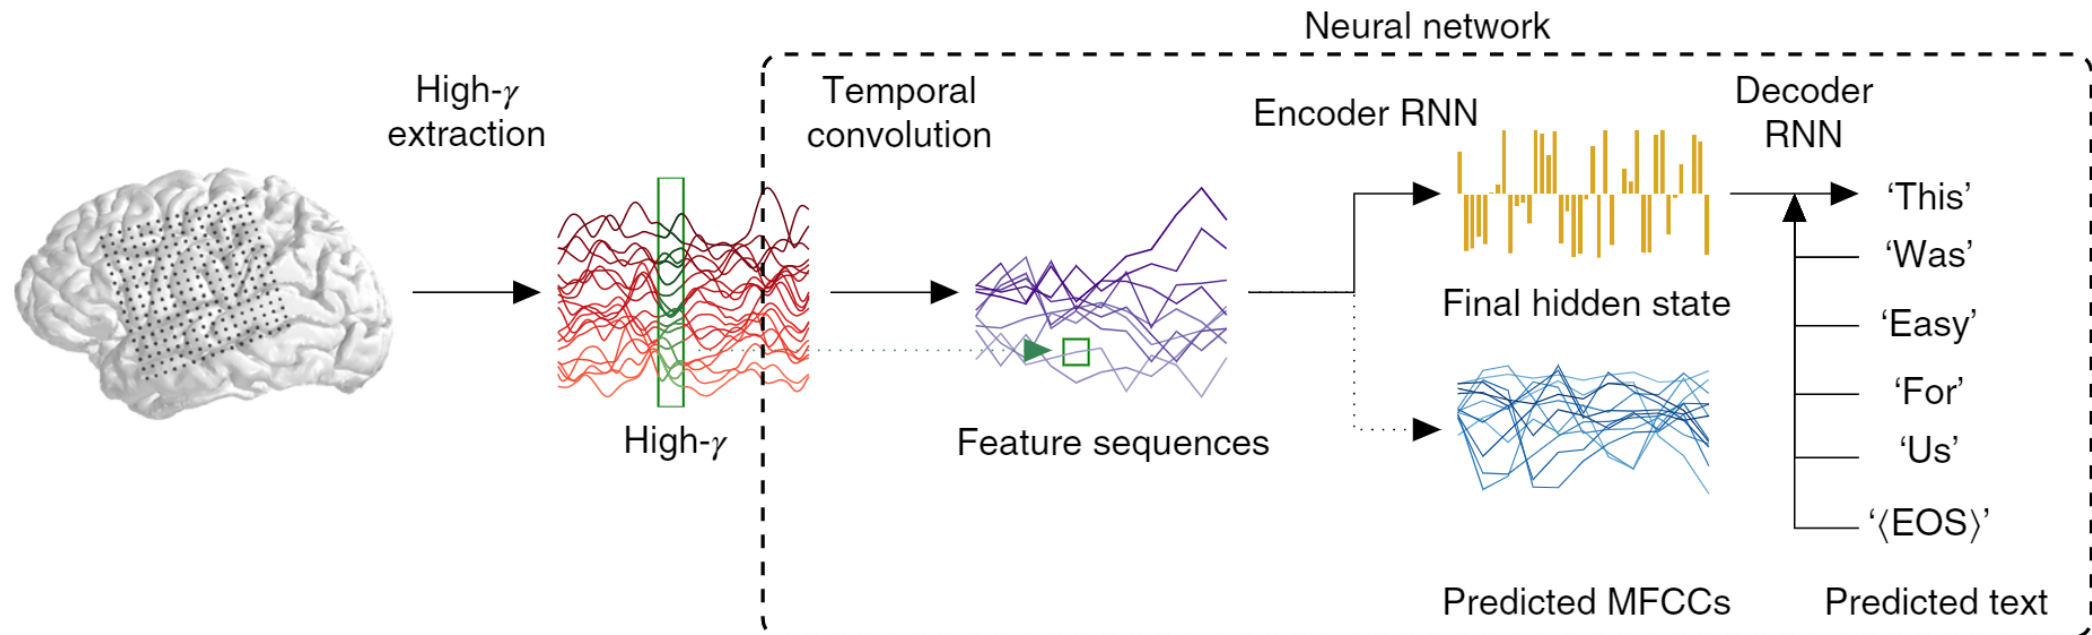


Application of Deep Neural Nets (1/2)

Brain decoding using recurrent neural net

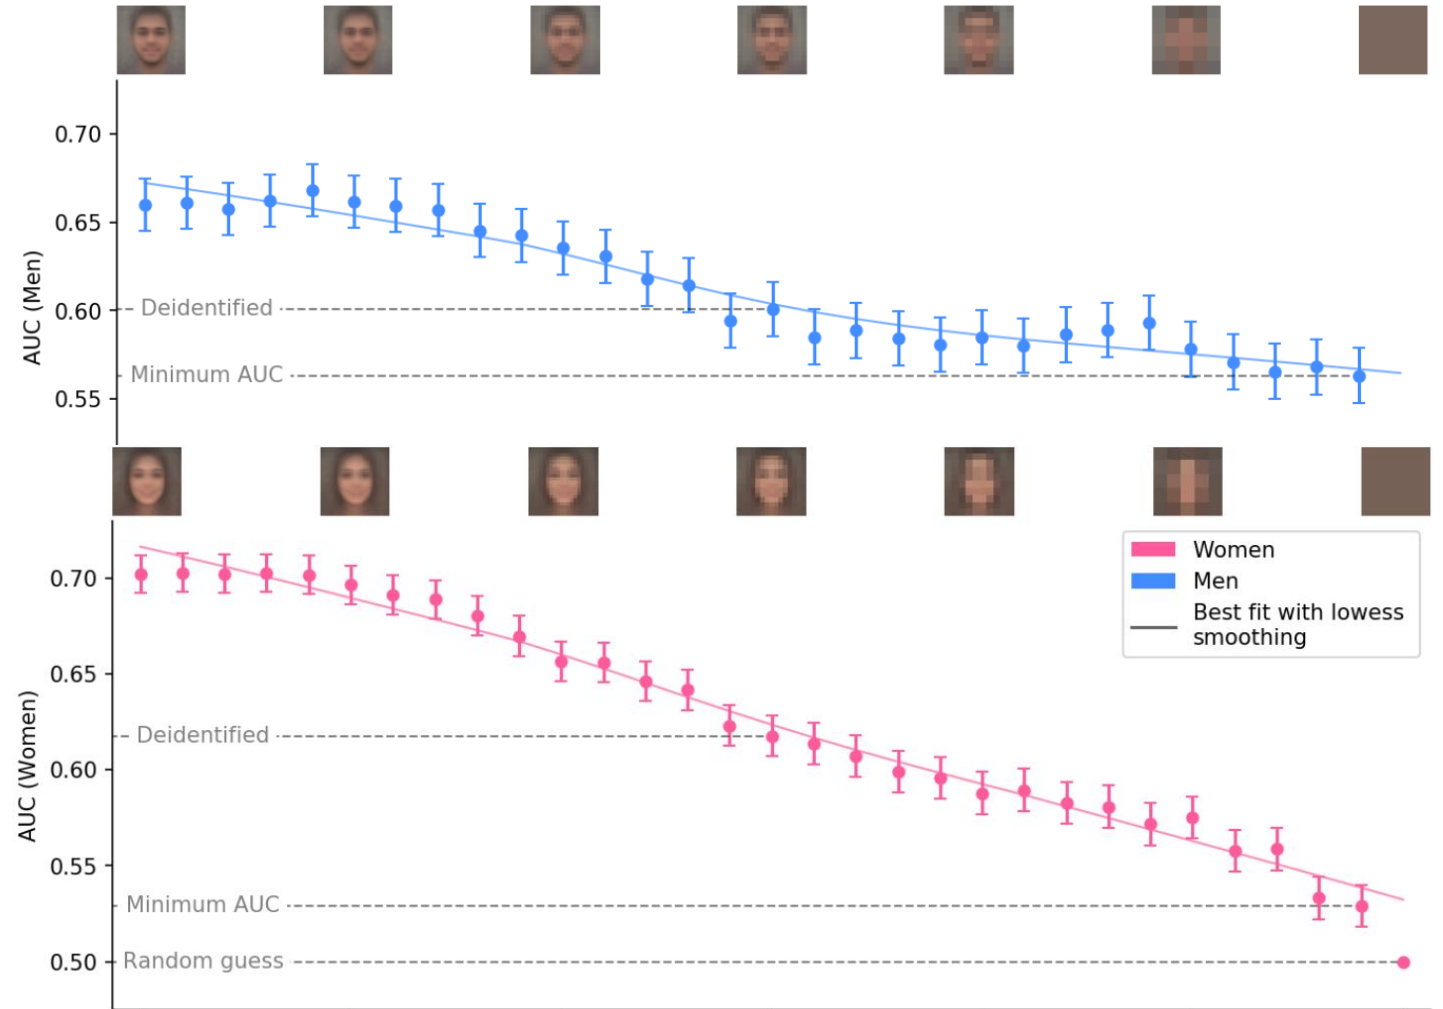
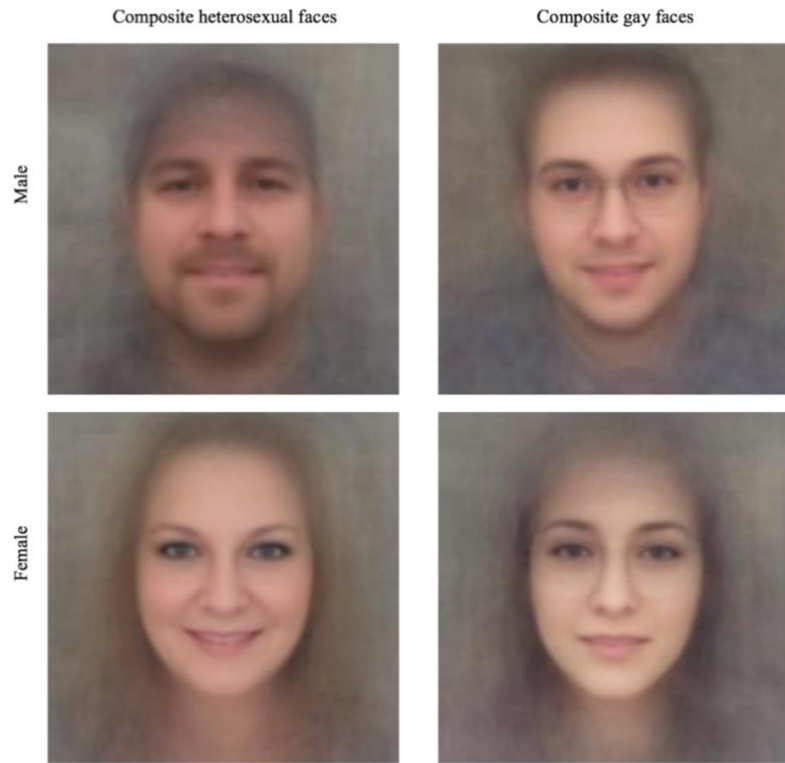
Machine translation of cortical activity to text with an encoder-decoder framework

Joseph G. Makin^{1,2}, David A. Moses^{1,2} and Edward F. Chang^{1,2}



Application of Deep Neural Nets (2/2)

Predicting sexuality using convolutional neural net



Topics for today

Computations of Neural Networks

Why do neural nets work?

(Deep) Learning in Neural Networks

How do neural nets learn?

Explainable AI

What do neural nets learn?



Topics for today

Computations of Neural Networks

Why do neural nets work?

(Deep) Learning in Neural Networks

How do neural nets learn?

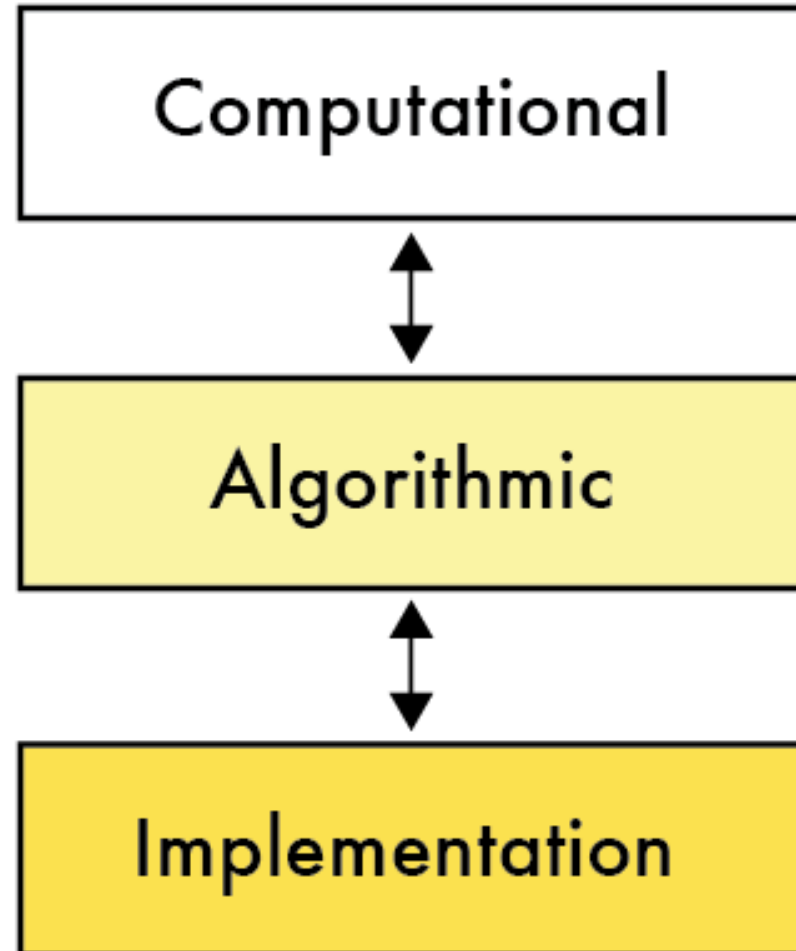
Explainable AI & Causal ML

What do neural nets learn?



Analysis of a Cognitive System

David Marr's 3 levels of analysis:



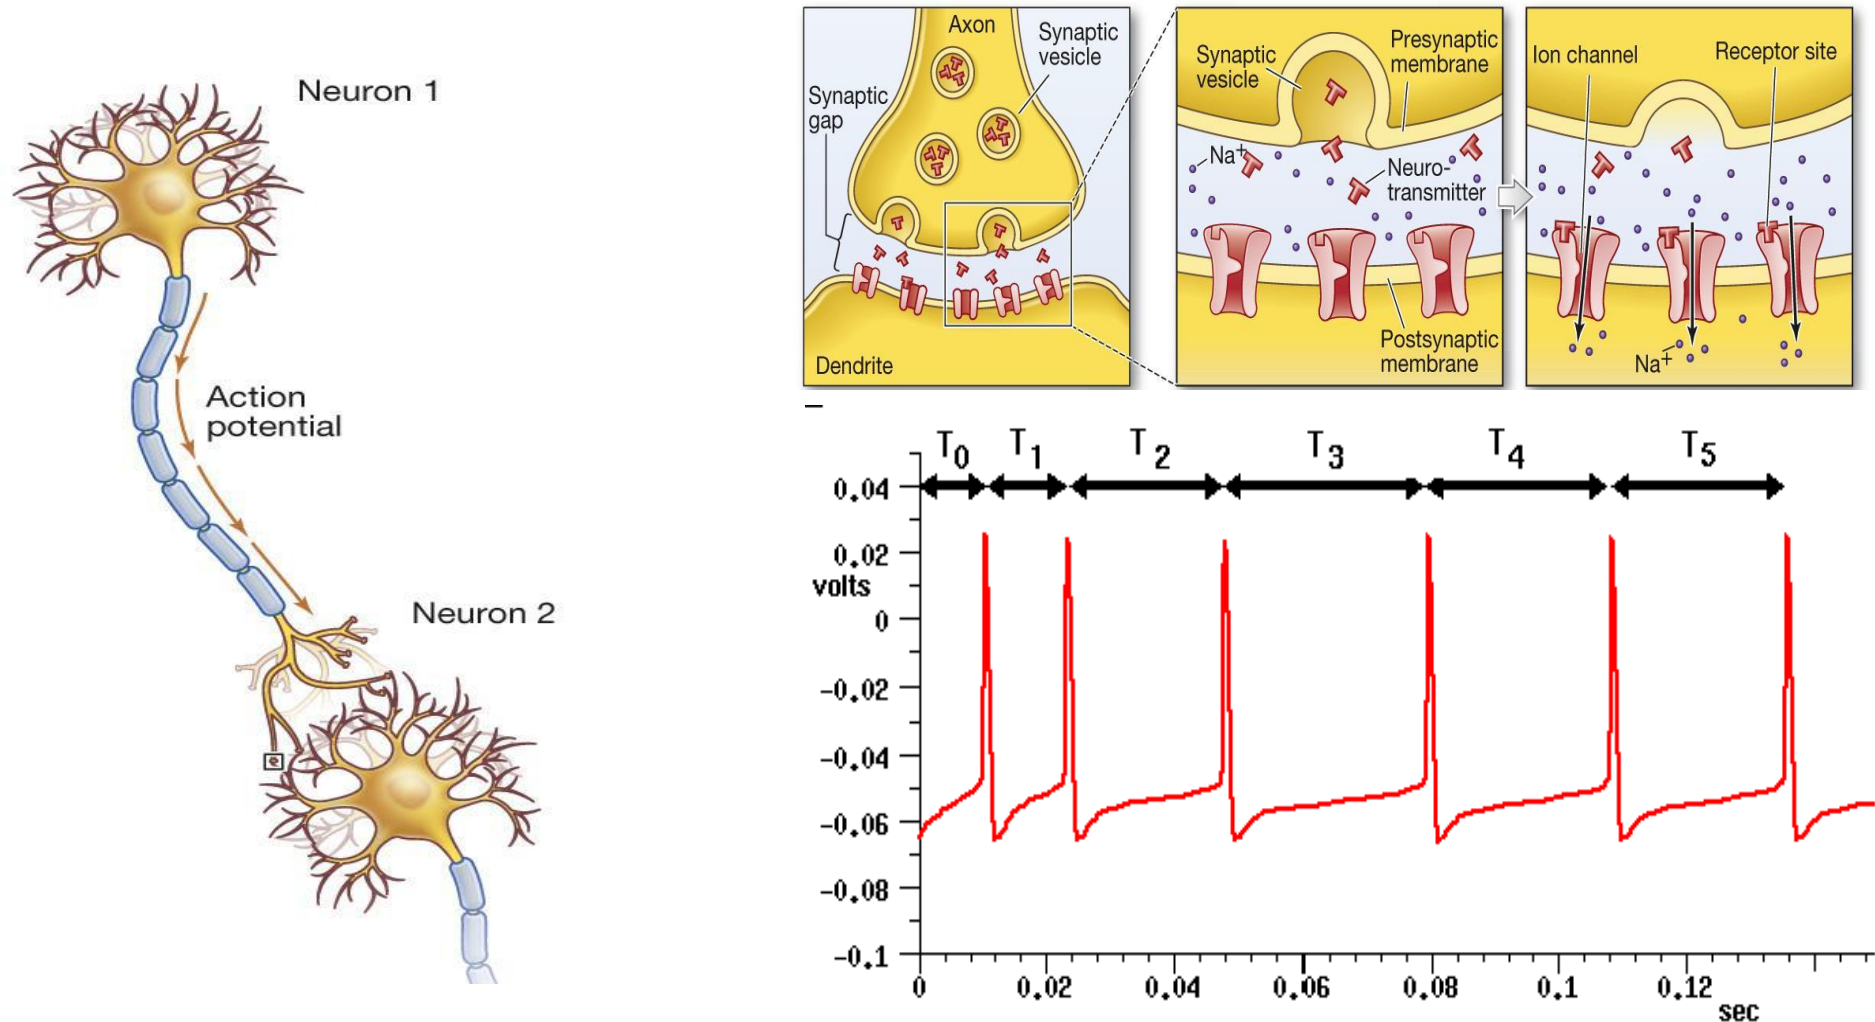
- What's the computational problem?

- How to solve it algorithmically?

- How to implement the algorithm?



Implementation (1/5): Possible States

A neuron has two states: 0 (resting) vs. 1 (firing)



Implementation (2/5): Addition

Suppose the neuron Z has a **low** threshold





X	Y	Z
0	0	0
0	1	1
1	0	1
1	1	1

$$Z = X + Y \text{ (OR)}$$

Implementation (3/5): Multiplication

Suppose the neuron Z has a **high** threshold





X	Y	Z
0	0	0
0	1	0
1	0	0
1	1	1

$$Z = X * Y \text{ (AND)}$$

Implementation (4/5): Division

From excitatory to **inhibitory** connections:



X	Z
0	1
1	0



$$Z = 1 - X \text{ (NOT)}$$

Implementation (5/5): Division


Inhibitions can lead to **subtraction** or **division**:

$$8 \rightarrow 4 \rightarrow 2 \rightarrow 1$$

$/2 \quad /2 \quad /2$




■ E
■ E + I



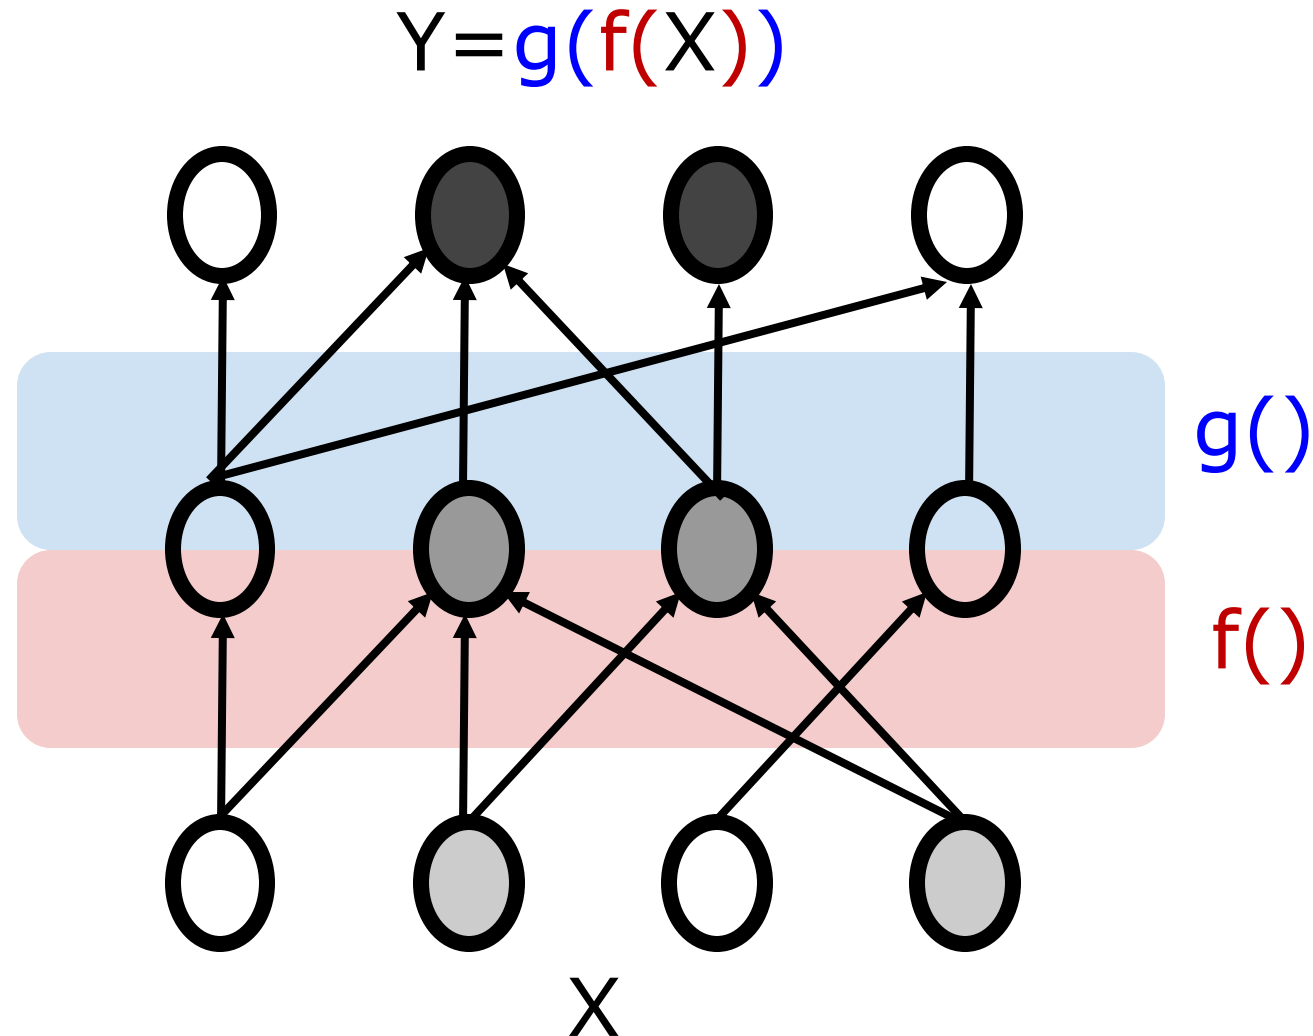
$$8 \rightarrow 6 \rightarrow 4 \rightarrow 2$$

$-2 \quad -2 \quad -2$




Algorithm: Arithmetic Combinations

A neural network=A series of continuous transformations



Computational Problem: Recognition






Here is a dog detector where W encodes a dog template:



Decision Boundary: $w_2 X_2 = -w_1 X_1 - w_0 \Rightarrow w_2 X_2 + w_1 X_1 + w_0 1 = 0$
 $(w_2, w_1, w_0) \cdot (X_2, X_1, 1) = 0$


Revisiting the ROC curve

A ML can transform features to solve its problem



Universal Approximation Theorem

A 3-layer net can approximate any continuous function




$$F(\mathbf{x}) = \sum_{i=1}^N v_i \varphi \left(\underbrace{w_i^T \mathbf{x}_i}_{y_j} - b_i \right)$$

as an approximate realization of the function f where f is independent of φ ;
that is,

$$|F(x) - f(x)| < \varepsilon$$

for all $x \in I_m$. In other words, functions of the form $F(x)$ are **dense** in $C(I_m)$.

Graph for $1/(1+\exp(-x))$, $1/(1+\exp((-5)*x))$, $1/(1+\exp((-5)*(x-2)))$



A 3-layer net, Taylor Series, & Fourier Transform are special cases of **Generalized Additive Models**!

Topics for today

Computations of Neural Networks

Why do neural nets work?

(Deep) Learning in Neural Networks

How do neural nets learn?

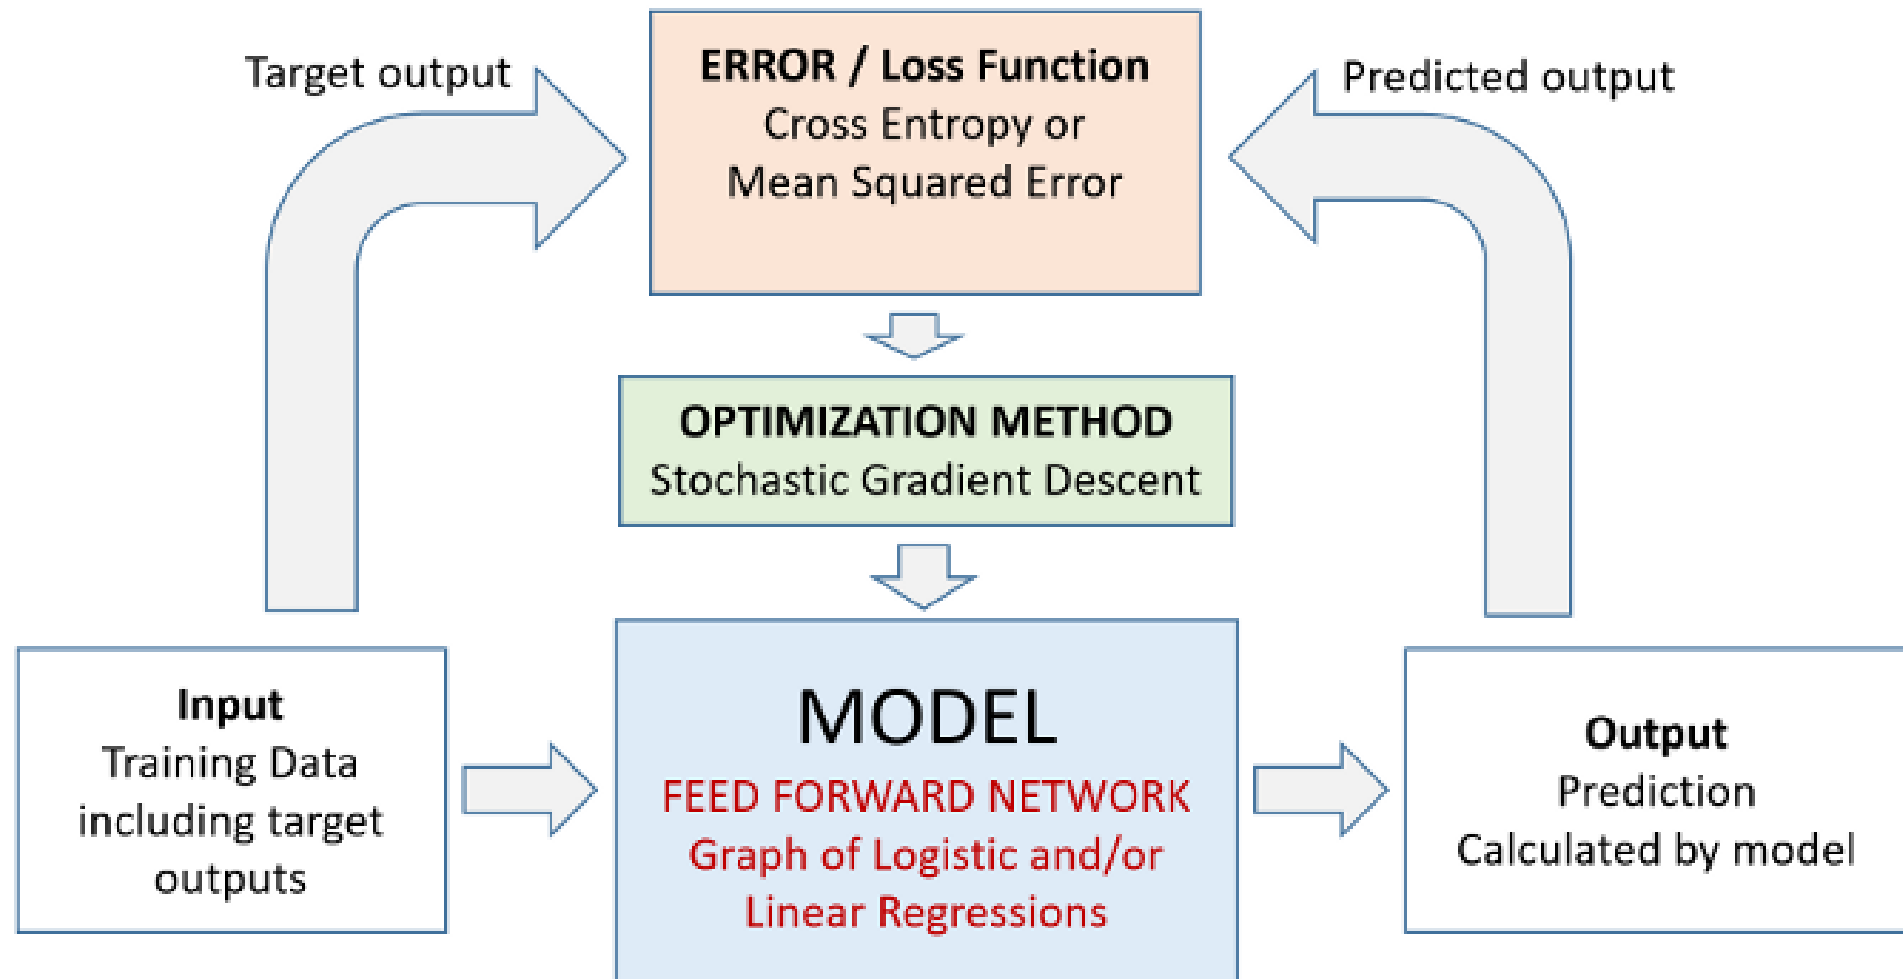
Explainable AI & Causal ML

What do neural nets learn?





Workflow of Supervised Learning

Adjusting model parameters to minimize prediction errors



Model (1/2): Frameworks

Keras is the easiest framework to build neural nets



Model (2/2): Keras Cheatsheet

Python For Data Science Cheat Sheet

Keras

Learn Python for data science interactively at [www.DataCamp.com](https://www.datacamp.com)



ERROR / Loss Function
Cross Entropy or
Mean Squared Error



OPTIMIZATION METHOD
Stochastic Gradient Descent



MODEL
FEED FORWARD NETWORK
Graph of Logistic and/or
Linear Regressions

Keras

Keras is a powerful and easy-to-use deep learning library for Theano and TensorFlow that provides a high-level neural networks API to develop and evaluate deep learning models.

A Basic Example

```
>>> import numpy as np
>>> from keras.models import Sequential
>>> from keras.layers import Dense
>>> data = np.random.random((1000,100))
>>> labels = np.random.randint(2,size=(1000,1))
>>> model = Sequential()
>>> model.add(Dense(32,
>>>                 activation='relu',
>>>                 input_dim=100))
>>> model.add(Dense(1, activation='sigmoid'))
>>> model.compile(optimizer='rmsprop',
>>>               loss='binary_crossentropy',
>>>               metrics=['accuracy'])
>>> model.fit(data, labels, epochs=10, batch_size=32)
>>> predictions = model.predict(data)
```

Data

Also see NumPy, Pandas & Scikit-Learn

Your data needs to be stored as NumPy arrays or as a list of NumPy arrays. Ideally, you split the data in training and test sets, for which you can also resort to the `train_test_split` module of `sklearn.cross_validation`.

Keras Data Sets

```
>>> from keras.datasets import boston_housing,
>>>                               mnist,
>>>                               cifar10,
>>>                               imdb
>>> (x_train,y_train),(x_test,y_test) = mnist.load_data()
>>> (x_train2,y_train2),(x_test2,y_test2) = boston_housing.load_data()
>>> (x_train3,y_train3),(x_test3,y_test3) = cifar10.load_data()
>>> (x_train4,y_train4),(x_test4,y_test4) = imdb.load_data(num_words=20000)
>>> num_classes = 10
```

Other

```
>>> from urllib.request import urlopen
>>> data = np.loadtxt(urlopen("http://archive.ics.uci.edu/
>>> ml/machine-learning-databases/pima-indians-diabetes/
>>> pima-indians-diabetes.data"),delimiter=",")
>>> X = data[:,0:8]
>>> y = data[:,8]
```

Preprocessing

Sequence Padding

```
>>> from keras.preprocessing import sequence
>>> x_train4 = sequence.pad_sequences(x_train4,maxlen=80)
>>> x_test4 = sequence.pad_sequences(x_test4,maxlen=80)
```

One-Hot Encoding

```
>>> from keras.utils import to_categorical
>>> Y_train = to_categorical(y_train, num_classes)
>>> Y_test = to_categorical(y_test, num_classes)
>>> Y_train3 = to_categorical(y_train3, num_classes)
>>> Y_test3 = to_categorical(y_test3, num_classes)
```

Model Architecture

Sequential Model

```
>>> from keras.models import Sequential
>>> model = Sequential()
>>> model2 = Sequential()
>>> model3 = Sequential()
```

Multilayer Perceptron (MLP)

Binary Classification

```
>>> from keras.layers import Dense
>>> model.add(Dense(12,
>>>                 input_dim=8,
>>>                 kernel_initializer='uniform',
>>>                 activation='relu'))
>>> model.add(Dense(8, kernel_initializer='uniform', activation='relu'))
>>> model.add(Dense(1, kernel_initializer='uniform', activation='sigmoid'))
```

Multi-Class Classification

```
>>> from keras.layers import Dropout
>>> model.add(Dense(512, activation='relu', input_shape=(784,)))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(512, activation='relu'))
>>> model.add(Dropout(0.2))
>>> model.add(Dense(10, activation='softmax'))
```

Regression

```
>>> model.add(Dense(64, activation='relu', input_dim=train_data.shape[1]))
>>> model.add(Dense(1))
```

Convolutional Neural Network (CNN)

```
>>> from keras.layers import Activation, Conv2D, MaxPooling2D, Flatten
>>> model2.add(Conv2D(32, (3,3), padding='same', input_shape=x_train.shape[1:]))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(32, (3,3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Conv2D(64, (3,3), padding='same'))
>>> model2.add(Activation('relu'))
>>> model2.add(Conv2D(64, (3,3)))
>>> model2.add(Activation('relu'))
>>> model2.add(MaxPooling2D(pool_size=(2,2)))
>>> model2.add(Dropout(0.25))
>>> model2.add(Flatten())
>>> model2.add(Dense(512))
>>> model2.add(Activation('relu'))
>>> model2.add(Dropout(0.5))
>>> model2.add(Dense(num_classes))
>>> model2.add(Activation('softmax'))
```

Recurrent Neural Network (RNN)

```
>>> from keras.layers import Embedding, LSTM
>>> model3.add(Embedding(20000,128))
>>> model3.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2))
>>> model3.add(Dense(1, activation='sigmoid'))
```

Also see NumPy & Scikit-Learn

Train and Test Sets

```
>>> from sklearn.model_selection import train_test_split
>>> X_train5, X_test5, y_train5, y_test5 = train_test_split(X,
>>>                                                         y,
>>>                                                         test_size=0.33,
>>>                                                         random_state=42)
```

Standardization/Normalization

```
>>> from sklearn.preprocessing import StandardScaler
>>> scaler = StandardScaler().fit(x_train2)
>>> standardized_X = scaler.transform(x_train2)
>>> standardized_X_test = scaler.transform(x_test2)
```

Inspect Model

```
>>> model.output_shape
>>> model.summary()
>>> model.get_config()
>>> model.get_weights()
```

Model output shape
Model summary representation
Model configuration
List all weight tensors in the model

Compile Model

MLP: Binary Classification

```
>>> model.compile(optimizer='adam',
>>>               loss='binary_crossentropy',
>>>               metrics=['accuracy'])
```

MLP: Multi-Class Classification

```
>>> model.compile(optimizer='rmsprop',
>>>               loss='categorical_crossentropy',
>>>               metrics=['accuracy'])
```

MLP: Regression

```
>>> model.compile(optimizer='rmsprop',
>>>               loss='mse',
>>>               metrics=['mae'])
```

Recurrent Neural Network

```
>>> model3.compile(loss='binary_crossentropy',
>>>                 optimizer='adam',
>>>                 metrics=['accuracy'])
```

Model Training

```
>>> model3.fit(x_train4,
>>>            y_train4,
>>>            batch_size=32,
>>>            epochs=15,
>>>            verbose=1,
>>>            validation_data=(x_test4, y_test4))
```

Evaluate Your Model's Performance

```
>>> score = model3.evaluate(x_test,
>>>                          y_test,
>>>                          batch_size=32)
```

Prediction

```
>>> model3.predict(x_test4, batch_size=32)
>>> model3.predict_classes(x_test4, batch_size=32)
```

Save/ Reload Models

```
>>> from keras.models import load_model
>>> model3.save('model_file.h5')
>>> my_model = load_model('my_model.h5')
```

Model Fine-tuning

Optimization Parameters

```
>>> from keras.optimizers import RMSprop
>>> opt = RMSprop(lr=0.0001, decay=1e-6)
>>> model2.compile(loss='categorical_crossentropy',
>>>                 optimizer=opt,
>>>                 metrics=['accuracy'])
```

Early Stopping

```
>>> from keras.callbacks import EarlyStopping
>>> early_stopping_monitor = EarlyStopping(patience=2)
>>> model3.fit(x_train4,
>>>            y_train4,
>>>            batch_size=32,
>>>            epochs=15,
>>>            validation_data=(x_test4, y_test4),
>>>            callbacks=[early_stopping_monitor])
```

DataCamp

Learn Python for Data Science interactively



Error/Loss Functions

Adjusting model parameters to minimize prediction errors

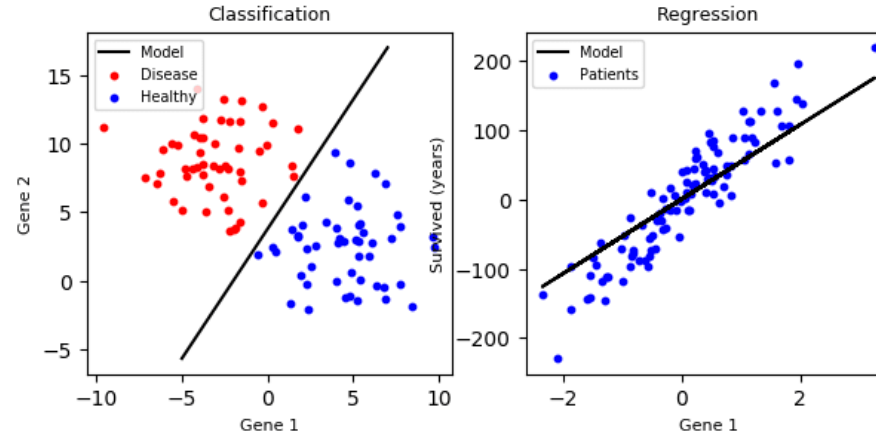
ERROR / Loss Function
Cross Entropy or
Mean Squared Error





OPTIMIZATION METHOD
Stochastic Gradient Descent



MODEL
FEED FORWARD NETWORK
Graph of Logistic and/or
Linear Regressions



Graph for $-\ln(x)$



$\hat{\mathbf{y}}$ \mathbf{y}

$$D(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_j y_j \ln \hat{y}_j$$

$\begin{bmatrix} 0.1 \\ 0.5 \\ 0.4 \end{bmatrix}$ $\begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$

When $\mathbf{Y}_{\text{pred}} = [0.1, 1, 0.4]$,
 $CE = -\ln(1) = 0$ but

$$MSE = \sqrt{0.1^2 + 0.4^2} \neq 0$$

Optimization (1/2): Gradient Descent

Using iterative methods to minimize error/loss functions

ERROR / Loss Function
Cross Entropy or
Mean Squared Error




OPTIMIZATION METHOD
Stochastic Gradient Descent




MODEL
FEED FORWARD NETWORK
Graph of Logistic and/or
Linear Regressions

$$\frac{\partial L(\vec{W})}{\partial \vec{W}} = 0 \xrightarrow[\text{over time}]{L \text{ is changing}} \vec{W}^{new} = \vec{W}^{old} - \alpha \frac{\partial L(\vec{W})}{\partial \vec{W}}$$

Convergence



Divergence



Optimization (2/2): Repetitions

1 batch/iteration = 1 adjustment of model parameters

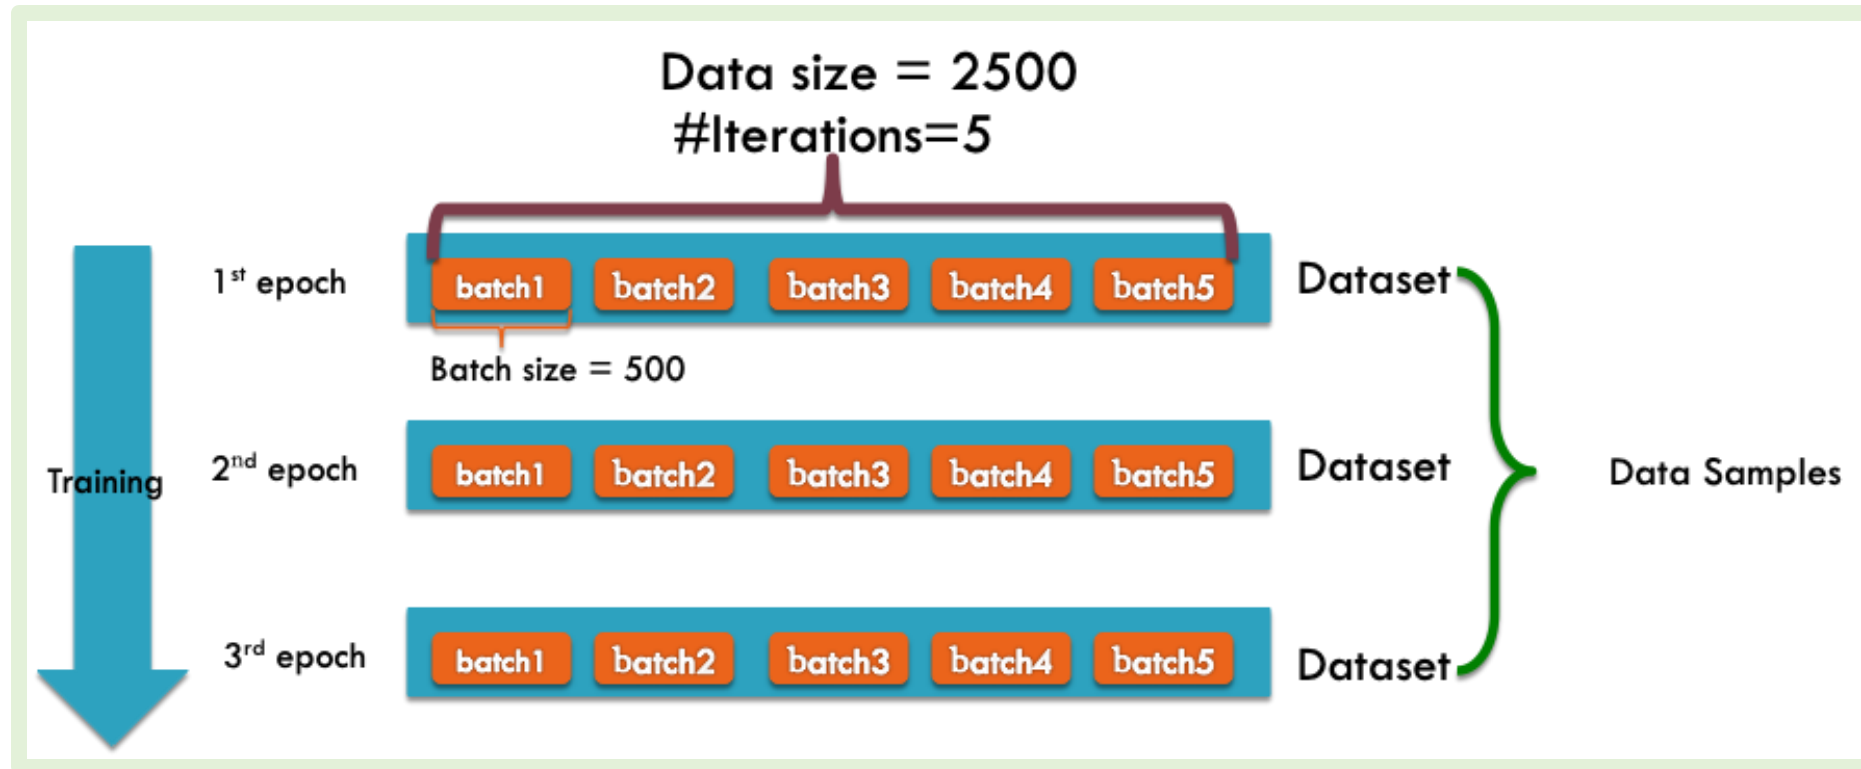
ERROR / Loss Function
Cross Entropy or
Mean Squared Error



OPTIMIZATION METHOD
Stochastic Gradient Descent



MODEL
FEED FORWARD NETWORK
Graph of Logistic and/or
Linear Regressions

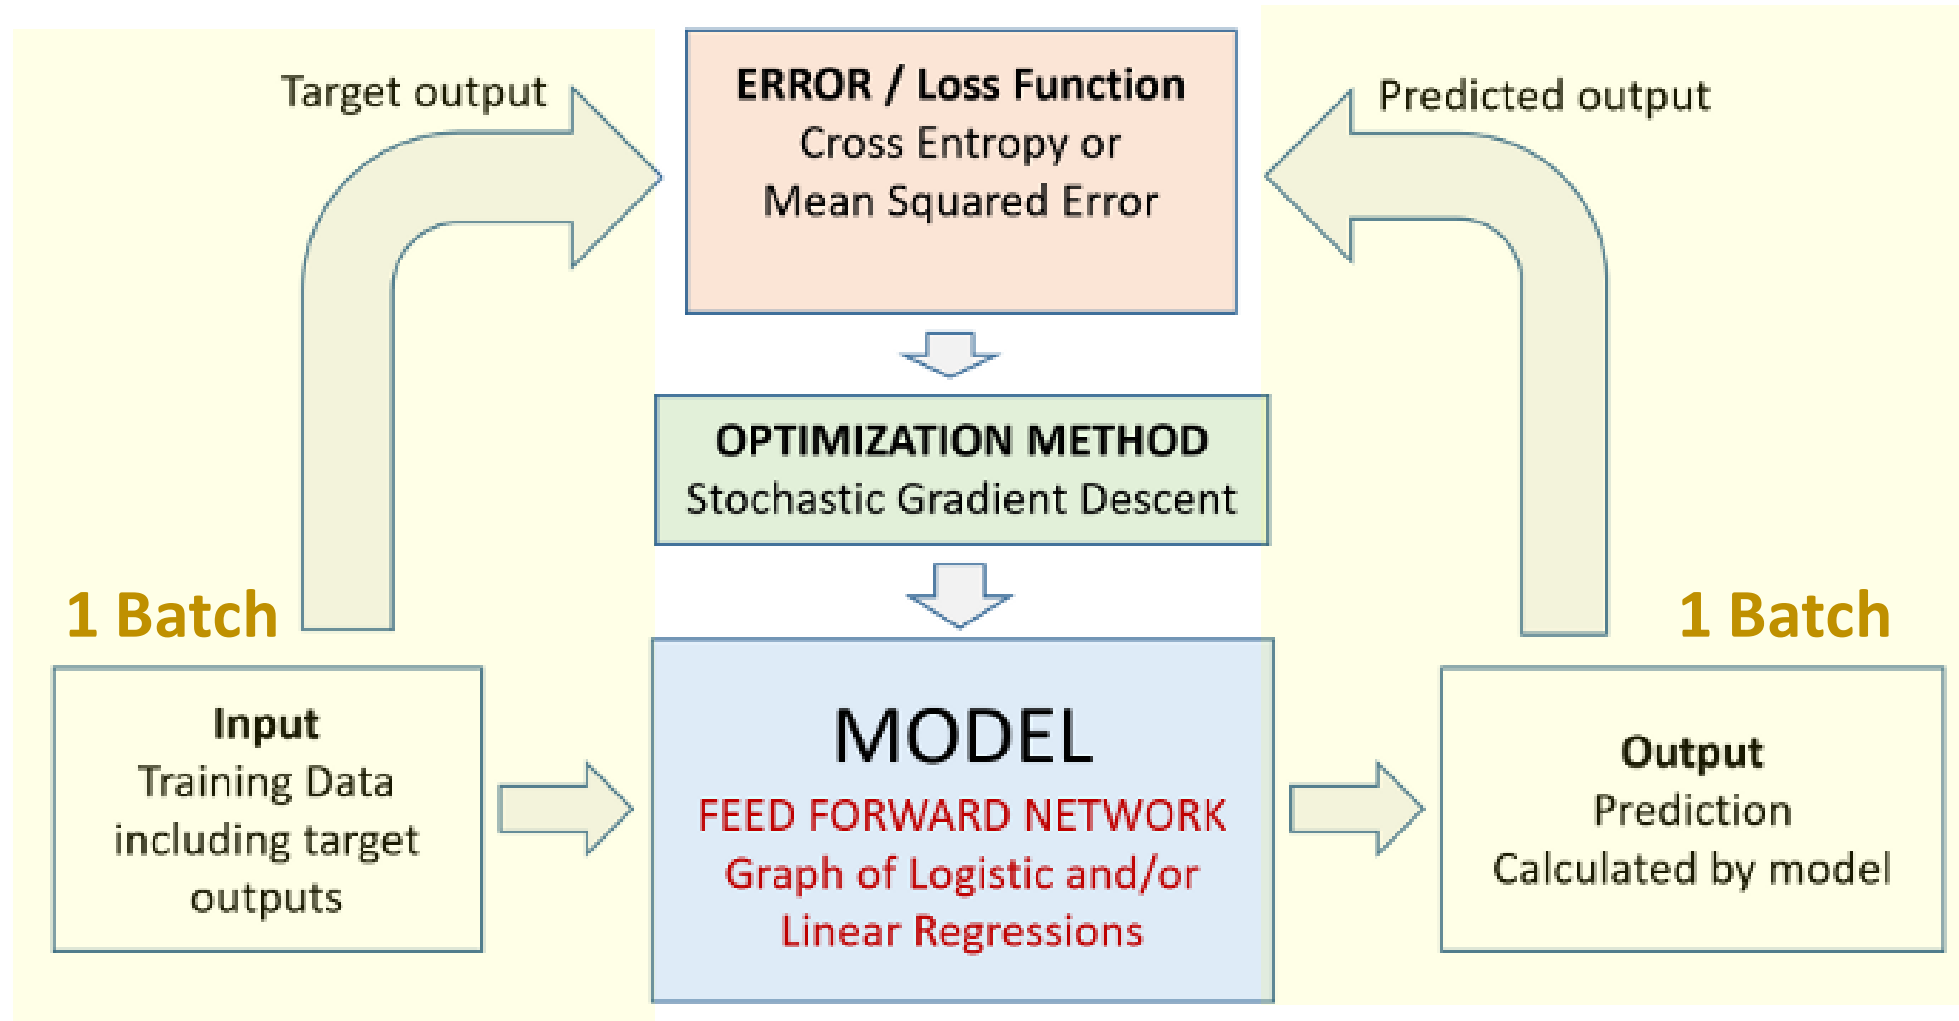


In each batch, L is different:

$$\vec{W}^{new} = \vec{W}^{old} - \alpha \frac{\partial L(\vec{W})}{\partial \vec{W}}$$

Workflow of Supervised Learning

Adjusting model parameters to minimize prediction errors



Topics for today

Computations of Neural Networks

Why do neural nets work?

(Deep) Learning in Neural Networks

How do neural nets learn?

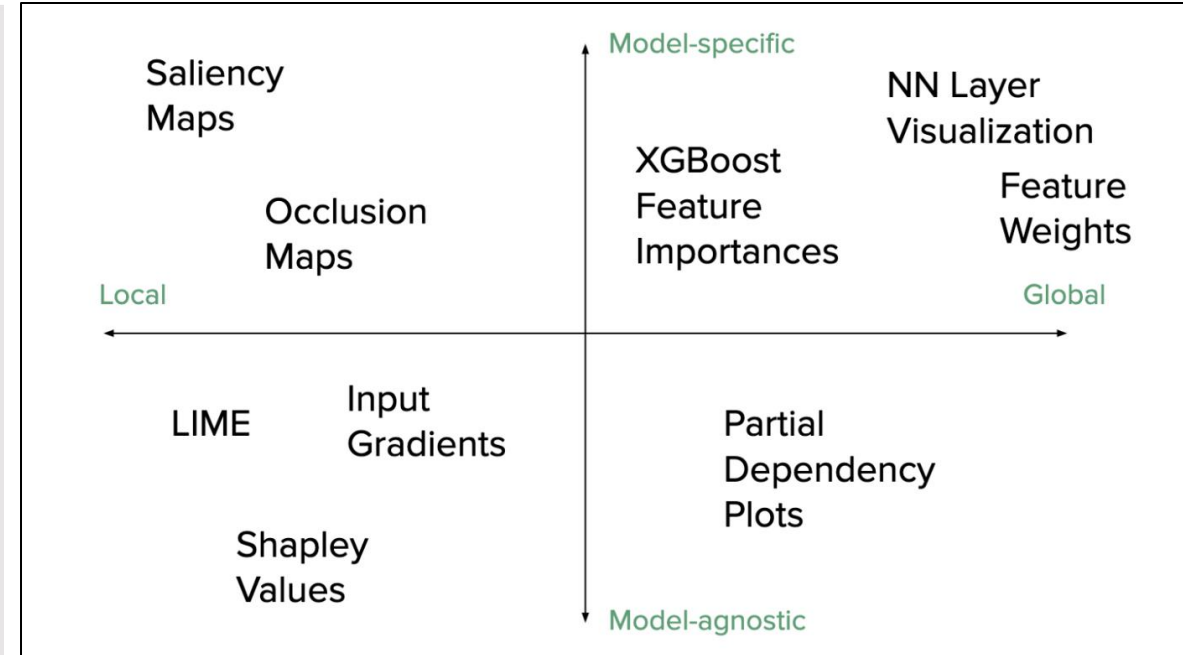
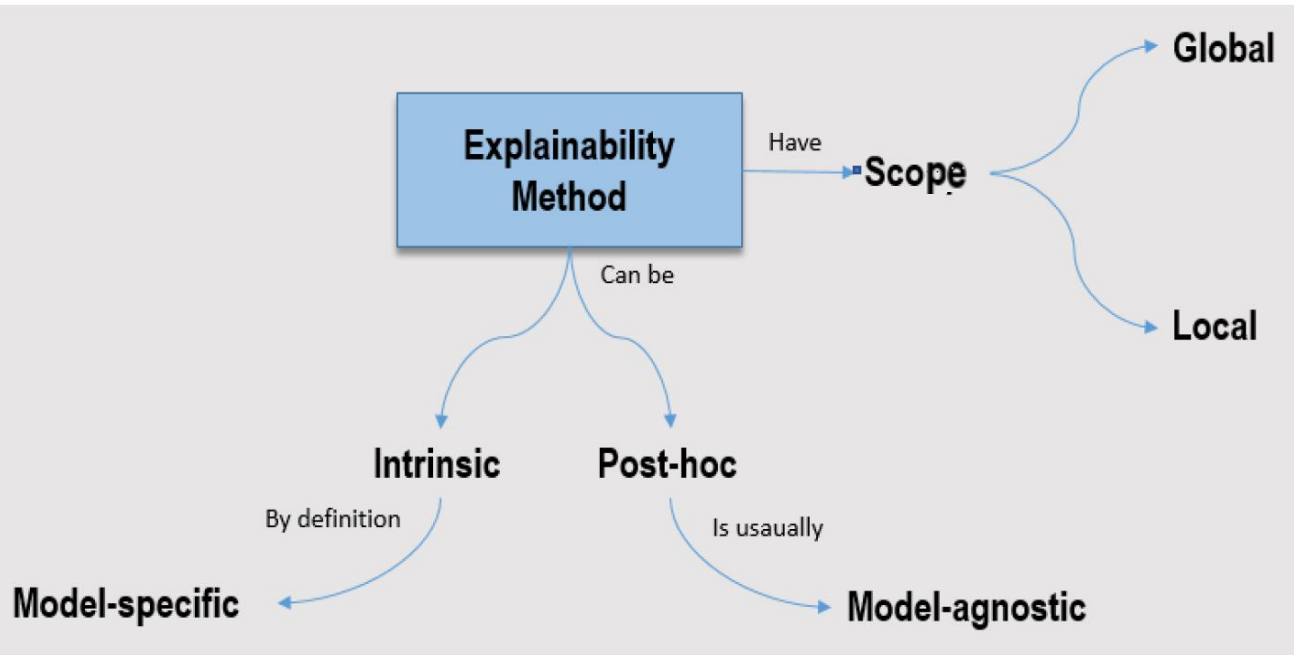
Explainable AI & Causal ML

What do neural nets learn?





eXplainable AI (XAI) methods

Model-agnostic: Models are accessible but not trainable

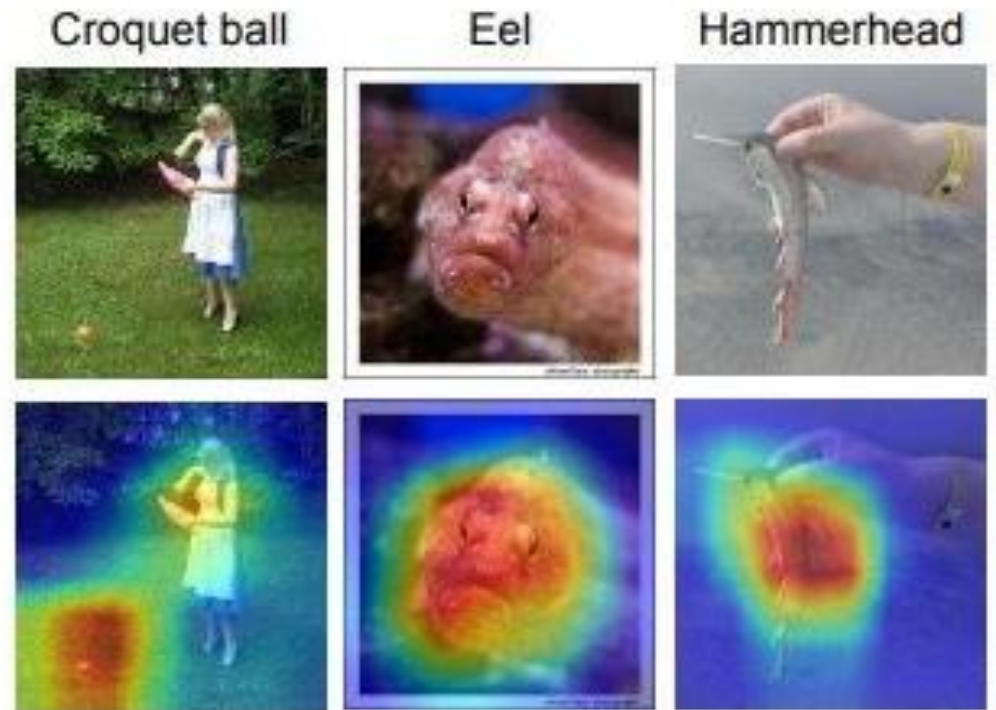


Cognitivist XAI: Model-specific

For example, simply adding a learnable attention layer

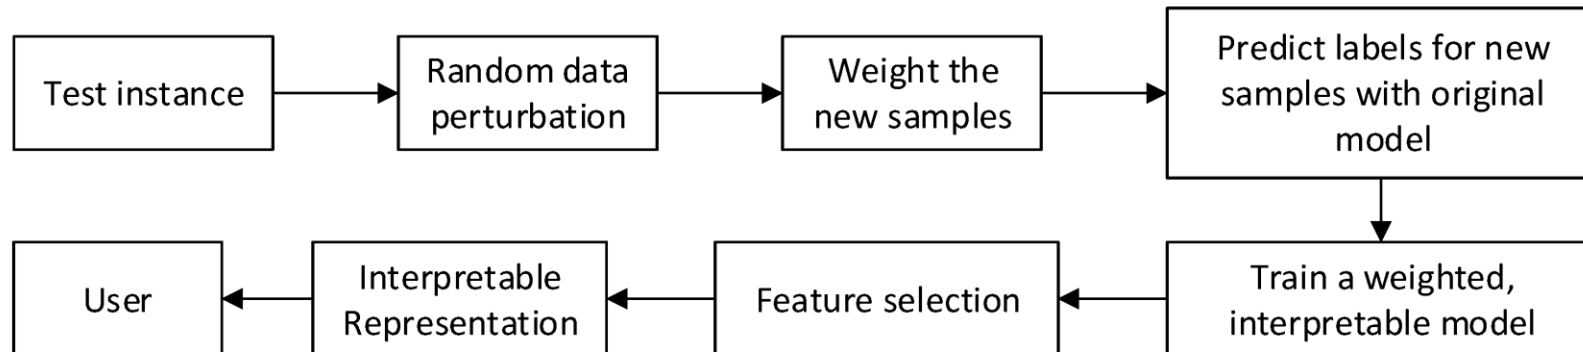




Input
image




Behaviorist XAI: Model-agnostic (1/2)

LIME=Local Interpretable Model-Agnostic Explanation



Behaviorist XAI: Model-agnostic (2/2)

LIME is a special case of SHapley Additive exPlanations (SHAP), which estimate **feature importance** better than **feature lesion**.




Ground Truth: $Y = X_1 + X_2 + X_1X_2$

Intact: $Y = \text{ModelA}(X_1, X_2)$

Lesioned: $Y = \text{ModelB}(X_2) = X_2$

Intact-Lesioned = $X_1 + X_1X_2$




But this difference is NOT solely contributed by X_1 !

Causal Machine Learning: Causality→AI

XAI may reveal that your AI is not learning real causality

An "accurate" classifier of
bile duct stone (膽管結石):



+



=




Panda
(60% confidence)

Adversarial
Perturbation

Gibbon
(99% confidence)

So what?

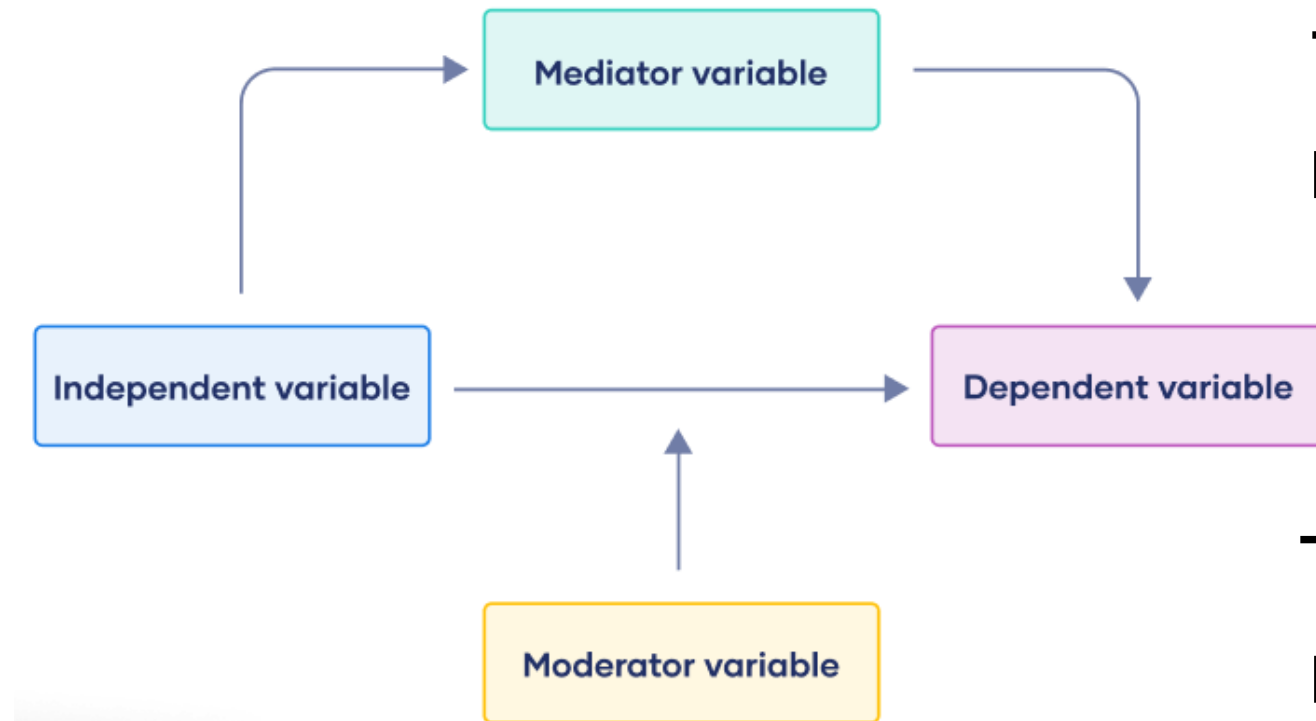
**Your model can't
generalize to
other datasets.**



Causal ML: AI → Causality (1/2)

The feature importance of X decreases after M added:

$$Y=f(X) \rightarrow Y=\mathbf{f}(X, M)=\beta_X X+\beta_M M+\beta_0$$



The linear function \mathbf{f} can be replaced by a nonlinear ML!

The linear function \mathbf{g} can be replaced by a nonlinear ML!


More accurate predictions after M added:

$$Y=g(X) \rightarrow Y=\mathbf{g}(X, M)=\beta_X X+\beta_M M+\beta_{MX} MX+\beta_0$$

Causal ML: AI → Causality (2/2)

Double ML uses ML to factor out confounding associations:

Confounding
Association



$$Y = f(X, W)$$

Step 1:

$$Y = g(W) + Y_R$$
$$X = h(W) + X_R$$

Step 2:

$$Y_R = \beta X_R + \varepsilon$$

Topics for today

Computations of Neural Networks

Why do neural nets work?

(Deep) Learning in Neural Networks

How do neural nets learn?

Explainable AI & Causal ML

What do neural nets learn?



GAME Over

